

Running KLEE on GNU coreutils

—

Report

Valentin Huber

Institute of Applied Information Technology
Zürich University of Applied Sciences ZHAW
contact@valentinhuber.me

February 12, 2024

Abstract

In 2008, Cadar *et al.* presented KLEE: a symbolic execution based fuzzer. In their report [1], they evaluate the performance of KLEE by running it against 89 programs from the GNU coreutils suite. In this work, I attempted to reproduce these results along with further experiments. Specifically — after a short introduction to KLEE and symbolic execution based fuzzing in general — I examined the statistical variance introduced by the inherent randomness of KLEE’s design, along with how changing the timeout passed to KLEE changes the measured outputs. I further ran the same experiments on two additional more recent versions of coreutils to see how KLEE performs on current software and to get a sense of how software develops over time. I then discuss how the results obtained in these experiments are analyzed and visualized. Finally, I reflect on this project, including a presentation of the produced artifacts and ideas for future work.

Contents

1	Introduction	1
2	Background	1
2.1	A Primer on Symbolic Execution	1
2.2	Symbolic Execution in Practice	2
3	Reproducing the Original Paper	2
3.1	Project Setup	2
3.2	Naïve Approach	3
3.3	Using an Old Version of Ubuntu	3
3.4	Generating LLVM Bytecode Files	3
3.5	Coverage Data Gathering	5
3.6	Preparing KLEE	5
3.7	Running KLEE	6
3.8	Extracting Human-Readable Coverage Data	7
3.9	Gathered Metrics	7
4	Comparing Runs	7
4.1	Comparison to the Original Paper	9
4.2	Influence of Timeout	9
4.2.1	Changes in Coverage	9
4.2.2	Changes in Number of Errors	12
5	Testing More Recent Versions of coreutils	13
5.1	Differences in Testing Setup	13
5.2	Findings	14
5.2.1	Changes in Coverage	14
5.2.2	Changes in Number of Errors	14
6	Discussion	16
6.1	Research Questions	16
6.2	Produced Artifacts	16
6.3	Future Work	17
	Bibliography	18
	Appendix	19
1	Building coreutils 6.10 on KLEE’s Docker image	19
1.1	Error	19
1.2	freadahead.c	19
2	Spread Plots of All Measurements	21
2.1	By Timeout	21
2.2	By Version	26

1 Introduction

KLEE [2] is an open source, symbolic execution based, advanced fuzzing platform. It was introduced in the seminal paper titled “KLEE: Unassisted and Automatic Generation of High-Coverage Tests for Complex Systems Programs” in 2008. In their article, Cadar *et al.* present their work and evaluate it on a diverse set of programs. The most prominent of those is the GNU coreutils suite, in which ten fatal errors were found.

Ever since then, KLEE has not only matured as a fuzzer, it has also been used extensively as a platform for other researchers to build on top of, as I have discovered in [3]. As an introduction to the practical side of fuzzing, I attempted to answer the following questions about KLEE:

1. Reproducing the original paper (see Section 3)
 - (a) Can the current version of KLEE be run on coreutils version 6.10 (as tested in the original paper)?
 - (b) Can the same metrics as measured in the original paper still be measured?
2. Examining the statistical distribution of results over different fuzzing times (see Section 4)
 - (a) How does the non-determinism in KLEE influence the variance in the results between different test runs?
 - (b) How do the measured metrics compare to what was published 15 years ago?
 - (c) How do different testing timeouts influence results?
3. Testing more recent versions of coreutils (see Section 5)
 - (a) What needs to change in the test setup to test more recent versions of coreutils?
 - (b) How do the results from testing different versions of coreutils differ?

All experiments were run on a virtualized server with the following specs: AMD EPYC 7713 64C 225W 2.0GHz Processor, 1 TiB RAM, 2x 25GiB/s Ethernet. The hardware was graciously provided by the Institute of Applied Information Technology at Zürich University of Applied Sciences [4].

2 Background

What follows is a short explanation of the application of symbolic execution in fuzzing. For more extensive background, I refer to some of my previous work [3], [5].

Remember that KLEE is an open-source, symbolic execution based fuzzer. It takes LLVM bytecode from the program under test (PUT) as its input, runs its analysis on it. KLEE then outputs some statistics about the run, inputs to the PUT that crash it, and inputs that, when executed, cover all branches KLEE has examined during its analysis.

2.1 A Primer on Symbolic Execution

KLEE is a fuzzer based on symbolic execution. This means that instead of executing a PUT with a concrete value, it instead runs through the instructions and maps relationships between data in memory (such as variables and user input) to mathematical formulas. So an instruction like `%result ← = add i32 %a, %b` would be mapped to the logical relationship $\phi_k = \phi_i + \phi_j$. Conditional jumps are mapped to conditions on these variables for both outcomes of the condition, so the instruction `%isPositive = icmp sgt i32 %result, 0` would be represented with $\phi_k > 0$ and $\phi_k \leq 0$ respectively.

The set of all conditions along a certain path through the PUT are called the *path condition*. It can be passed to a satisfiability modulo theories (SMT) solver (KLEE uses STP [6] as a default), which

returns values for all user inputs, such that the PUT is forced down the exact path represented by the path condition.

This is the major advantage of symbolic execution based fuzzing, as compared to ordinary fuzzing. By not using concrete values, but instead logical representations of user input, it essentially runs through the PUT with all possible user inputs *at the same time*. So if the solver returns that no inputs satisfy the passed formula, we have proven that such inputs simply do not exist. To be able to do this, it accepts the huge overhead of translating the code to formulas and then solving them.

2.2 Symbolic Execution in Practice

Symbolic execution in fuzzing has several major challenges to overcome. I have previously discussed them in detail [3], but would like to give a short summary here:

- Environment interactions (such as file system interactions) in general are opaque to the fuzzer and cannot be mapped to logical formulas. KLEE deals with this by solving the path constraint before the instruction in question and then uses concrete values in the call. This abandons the claim on completeness symbolic execution typically has, but is often the only feasible way to still continue analysis.
- The second major challenge in symbolic execution is what is known as *path explosion*. Because the number of program states grows exponentially with the number of instructions, for all but the most simple programs it is not feasible to calculate the entire state space. KLEE deals with this by reducing the search space to actually executable instructions, using advanced data structures, and examining paths through the PUT consecutively, with a user-defined timeout. To maximize the state space and code coverage as quickly as possible, it alternates between two strategies for choosing the next input to evaluate: KLEE either chooses the input that promises to increase the coverage the most and a random input to prevent the execution from getting stuck in a certain subtree of the PUT.
- KLEE needs to model the entire memory of a process. This is straight-forward as long as variables are used directly but becomes a challenge when pointers are involved. This is especially true if the value of these pointers depends on user input, since this would require KLEE to model all possible addresses having all possible values, which instantly explodes the memory consumption and number of states to examine and is thus infeasible. KLEE deals with this by representing such pointer operation as array accesses where the accessed object is copied as often as necessary to model all possible results, including error states.
- As programs become more complex, the path constraints become increasingly long and solving them contributes more and more to the fuzzer's runtime. KLEE applies some advanced optimizations, like query splitting and more general optimizations, or a cache of previous results, which often solve supersets the query they are a solution to. Finally, KLEE defines a timeout, after which the solver is interrupted and analysis is continued at an other branch.

3 Reproducing the Original Paper

I'm basing my experiment setup on the original paper [1], the FAQs in the project's documentation [7] and the tutorial on testing coreutils version 6.11 [8].

3.1 Project Setup

KLEE is a complex system including complex dependencies such as the SMT solvers. The maintainers provide a Dockerfile and the corresponding Docker image. Using Docker as an intermediate form of virtualization adds a layer of indirection and a performance penalty. However, since I'm not necessarily

interested in maximizing performance in this project, but instead focus on comparing different setups, this is a tradeoff worth taking. Using Docker to evaluate fuzzers’ performance has been done before [9]. Finally, this makes complex build steps reproducible and acts as documentation.

3.2 Naïve Approach

When attempting to build coreutils 6.10 directly in the current version of KLEE’s Docker image, I ran into an issue: The Docker image is based on Ubuntu 22 (Jammy), and no longer is able to build coreutils 6.10 with the GNU Compiler (GCC). This is because coreutils’ build system attempts to detect what system it is running on, and the variable the detection is based on is no longer defined. Specifically, in `freadahead.c` the following check is performed:

```
25 #if defined _IO_ferror_unlocked /* GNU libc, BeOS */
```

The error message and the full `freadahead.c` can be found in Appendix 1.

3.3 Using an Old Version of Ubuntu

One attempt to mitigate this issue would be to rewrite this check to allow the version of GCC installed on KLEE’s Docker image to compile coreutils 6.10. However, I opted to pursue a different avenue, because of two reasons:

1. Build systems are not my area of expertise and I do not know how many other issues would appear once the first was solved.
2. Changing code always adds risk of introducing additional software errors, which would distort my findings.

Therefore, I attempted to build the binaries on an old version of Ubuntu, and then move the binaries to KLEE’s Docker image. Specifically, I chose the latest LTS version which was available when version 6.10 of coreutils was current. This approach worked without any additional changes to the code nor the build system. The setup of the Docker image then used to build coreutils can be seen in Listing 1.

3.4 Generating LLVM Bytecode Files

Building binaries themselves is unfortunately not enough, since KLEE does not take pure binaries as its input, but instead requires LLVM bytecode. Compiling an ordinary `.c` file to LLVM can easily be done using `clang`. However, again, coreutils use a complex build system which means to just use `clang`, I’d have to deeply understand and modify it, with the drawbacks listed above.

Simply passing `clang` as the C compiler to the build system does not work, since the produced output is not a runnable binary, and the build system requires the compiler’s output to be executable.

Fortunately, there exists Whole Program LLVM (WLLVM) [10], a tool specifically designed to work with complex build systems while still producing LLVM bytecode as one of its outputs. This is achieved by injecting its compiler into the build system. The compiler creates executable binaries and additionally injects LLVM bytecode into a dedicated section of the object files. In a second step, these then get extracted and linked together to produce LLVM bytecode files.

Since I’m running WLLVM on an old version of Ubuntu, I was forced to use an old version of WLLVM as well, because newer versions require a version of python which is not available on Ubuntu 14.04. To create proper input files for KLEE, I added two options, to reduce warnings (`--build`) and to turn off premature optimizations according to the KLEE documentation (`CFLAGS`) [8].

The Dockerfile section to build the LLVM bytecode can be found in Listing 2.

LISTING 1: Dockerfile content to prepare a system for building coreutils 6.10

```

1 # =====
2 # base
3 # =====
4
5 FROM ubuntu:14.04 as klee-coreutils-base
6
7 # installing dependencies
8 RUN apt-get update \
9     && apt-get install -y \
10     wget \
11     build-essential
12
13 # downloading source code
14 RUN wget "http://ftp.gnu.org/gnu/coreutils/coreutils-6.10.tar.gz" \
15     && tar xf "coreutils-6.10.tar.gz" \
16     && mv "coreutils-6.10" coreutils
17
18 # modifying source code according to the documentation of the original experiment
19 RUN sed -i \
20     's/^#define INPUT_FILE_SIZE_GUESS (1024 \* 1024)$/#define INPUT_FILE_SIZE_GUESS
    ↪ 1024/g' \
21     coreutils/src/sort.c

```

LISTING 2: Dockerfile content to build coreutils to LLVM bytecode using WLLVM

```

23 # =====
24 # llvm
25 # =====
26
27 FROM klee-coreutils-base as klee-coreutils-llvm
28
29 # installing dependencies
30 RUN apt-get install -y \
31     clang \
32     llvm \
33     python3-pip
34
35 # Newer versions are no longer compatible with the latest python version available on
    ↪ Ubuntu 14.04
36 RUN pip3 install --upgrade -v "wllvm==1.1.5"
37
38 ENV LLVM_COMPILER clang
39 ENV CC wllvm
40
41 # compiling code to llvm bytecode (.bc)
42 WORKDIR /coreutils/obj-llvm
43 RUN ./configure \
44     --build x86_64-pc-linux-gnu \
45     --disable-nls \
46     LLVM_COMPILER=clang \
47     CC=wllvm \
48     CFLAGS="-O0 -D__NO_STRING_INLINES -D_FORTIFY_SOURCE=0 -U__OPTIMIZE__" \
49     && make \
50     && make -C src arch hostname
51
52 # extracting llvm bytecode from object files
53 WORKDIR /coreutils/obj-llvm/src
54 RUN find . -executable -type f | xargs -I '{}' extract-bc '{}'

```

LISTING 3: Dockerfile content to build coreutils instrumented to record coverage

```

56 # =====
57 # gcov
58 # =====
59
60 FROM klee-coreutils-base as klee-coreutils-gcov
61
62 # compiling code to binaries instrumented with gcov
63 WORKDIR /coreutils/obj-cov
64 RUN ./configure \
65     --build x86_64-pc-linux-gnu \
66     --disable-nls \
67     CFLAGS="-O2 -g -fprofile-arcs -ftest-coverage" \
68     && find .. -type f -name '*.c' -exec sed -i -E 's/\b_exit\(\)/exit(/g' {} + \
69     && make \
70     && make -C src arch hostname

```

3.5 Coverage Data Gathering

A simple way to compare what these experiments accomplish compared to the experiments documented in the original paper is to look at coverage data, specifically coverage as measured by `gcov`. To gather this information, one needs to compile the binaries using GCC, and tell the compiler to add coverage gathering instrumentation. Along with the binaries, a note document (`<executable-name>.gcno`) is created. When the binary is executed, the added instrumentation records which path through the code is taken and, together with information from the notes file, stores its results in a coverage data file (`<executable-name>.gcda`). This file can then be analyzed with `gcov` to get human-readable coverage data.

With this step however, I ran into the same issue as before: Recent versions of GCC no longer build coreutils 6.10. I adopted the same approach and used the same base image as described in Section 3.3. The Dockerfile excerpt with the build step can be found in Listing 3.

I made two changes compared to building the LLVM bytecode files, to increase the accuracy of the measurements:

- I replaced all calls to `_exit` with calls to `exit`, so that those instructions are also included in the measurements. This was done according to the instructions in the FAQ [7].
- The original paper mentions that coverage is measured only on executable lines of code. Specifically, Section 5.1 of the original paper says

“We measure size in terms of executable lines of code (ELOC) by counting the total number of executable lines in the final executable after global optimization, which eliminates uncalled functions and other dead code.” [1]

I am not sure how Cadar *et al.* calculated the executable lines of code, since this is not trivial. I did enable normal global optimization (`-O2`), but this may still result in a considerable underestimation of coverage.

3.6 Preparing KLEE

Finally, the bytecode files can be passed to KLEE for the actual fuzzing. To prepare KLEE’s Docker image, the environment and sandbox are prepared according to the documentation [7]. Then, the bytecode files from the step outlined in Section 3.4 and the binaries instrumented with `gcov` along with their notes files are copied to the analysis image. The analysis step itself is an involved process itself and is done by executing a shell script (`analyze.sh`). This step is explained in Section 3.7. The analysis

LISTING 4: Dockerfile content to prepare the fuzzing stage

```

84 # =====
85 # exec
86 # =====
87
88 FROM klee/klee AS klee-coreutils-exec
89
90 # setting up klee env
91 RUN wget "http://www.doc.ic.ac.uk/~cristic/klee/testing-env.sh" \
92     && env -i /bin/bash -c '(source testing-env.sh; env >test.env)' \
93     && wget "http://www.doc.ic.ac.uk/~cristic/klee/sandbox.tgz" \
94     && tar xzfv sandbox.tgz \
95     && mv sandbox.tgz /tmp \
96     && mv sandbox /tmp
97
98 # copying files from build stage
99 COPY --from=klee-coreutils-llvm --chown=klee /coreutils/ ./coreutils-llvm/
100 COPY --from=klee-coreutils-gcov --chown=klee /coreutils/ ./coreutils-gcov/
101
102 # copying run scripts
103 COPY analyze.sh ./
104
105 # setting default values for analyze script
106 # can be overridden using -e in docker run
107 ENV KLEE_MAX_TIME_MIN 60
108 ENV UTIL echo
109 ENV SKIP_KLEE_ANALYSIS ""
110
111 CMD bash ./analyze.sh \
112     --llvm-dir ./coreutils-llvm/obj-llvm/src \
113     --cov-dir ./coreutils-gcov/obj-cov/src \
114     --skip-klee-analysis "${SKIP_KLEE_ANALYSIS}" \
115     --klee-max-time "${KLEE_MAX_TIME_MIN}" \
116     --out-dir ./out \
117     "${UTIL}"

```

script is copied into the image and executed on container start. To allow passing certain settings to the analysis step, environment variables are used, which can be set in the `docker run` command.

The Dockerfile excerpt for this step can be found in Listing 4.

3.7 Running KLEE

When starting the Docker image built with the steps outlined before, a shell script is executed. This script handles the evaluation settings, input and output files, and collects metrics. Specifically, the following steps are performed:

1. The input including the passed settings are parsed. The script allows setting input and output directories (`--llvm-dir`, `--cov-dir`, `--out-dir`), KLEE's timeout (`--klee-max-time`), and skipping the fuzzing step (`--skip-klee-analysis`). The latter allows gathering additional metrics without based on the output from a previous fuzzing run without having to perform additional, computationally expensive analysis.
2. To run KLEE, the analyst is required to pass arguments setting the size and number of inputs and input files to be tested. For most coreutils, this is the same, however (as mentioned in Section 5.2 of the original paper [1] and explained in the FAQs [7]) some utils need different settings to achieve a decent coverage. The analysis script assembles the command to run KLEE, including the constant settings, util-dependant settings, and the timeout set in the script arguments.

3. Then, the actual fuzzing is performed.
4. KLEE’s output is examined in a few ways:
 - (a) For each found error, human readable outputs are created using `ktest-tool`.
 - (b) `klee-stats` is invoked to export metrics collected by KLEE.
 - (c) All test cases generated by KLEE are used as input for `klee-replay` pointed at the binary instrumented with coverage gathering instructions. This ensures that each instruction analyzed by KLEE during its fuzzing is also executed and thus recorded in the coverage results. Since the instrumented binaries were not compiled on the same system as they are executed on, the environment variables `GCOV_PREFIX` and `GCOV_PREFIX_STRIP` need to be set appropriately.
5. Finally, certain large output text files are compressed to minimize disk usage.

3.8 Extracting Human-Readable Coverage Data

As a last step, the output of the binaries instrumented to gather coverage metrics needs to be fed back into `gcov`. Unfortunately, the format of these output files changed at some point and the version of `gcov` installed in KLEE’s Docker image is no longer able to read them. They are therefore fed back into the Docker image that created them, where an obviously compatible version of `gcov` is available.

3.9 Gathered Metrics

In each run on each util of the experiment as discussed above, the following metrics were collected:

- The util name including the version and the run
- The timeout passed to KLEE
- The number of errors as reported by KLEE by type according to the file extension:
 - `ptr` errors, e.g. invalid pointers, null page accesses, out of bound pointers
 - `exec` errors, which appear on illegal instructions and external calls with a symbolic `errno` call
 - `model` errors, which appears when a symbolic size is concretized
 - `solver` errors, which are query timeouts
 - `abort` errors
 - The total number of errors
- Instruction and branch coverage¹ as reported by KLEE²
- Coverage as measured by `gcov`

4 Comparing Runs

In the following sections, only a small subset of all produced plots are discussed. Summary plots showing the spread and empirical cumulative distribution function (ECDF) of all measurements across timeouts and versions, and plots showing the results of individual utils are available in the repository of this experiment [11].

Figure 1 shows the spread of the measurements between four runs of the experiments with the same parameters. The non-deterministic nature of the results is due to the inherent non-determinism

¹KLEE includes library code in its coverage numbers, and thus reports significantly lower coverage compared to `gcov`.

²KLEE reports a set of additional metrics like time spent in the solver, number of instructions analyzed and number of cache hits, but these were not examined further in this paper.

of the fuzzing process in KLEE (refer to Section 2.2). Specifically, the standard deviations for code coverage by branch (according to KLEE), instruction (according to KLEE) and line (according to `gcov`) are 0.37%, 0.48%, and 1.56% respectively.




FIGURE 1: Spread of values normalized to the mean by util for coreutils 6.10 and a timeout of 60 minutes

Purely looking at the number of errors in each category is a very broad and imprecise way of measuring results. However, comparing and deduplicating errors found by a fuzzer is an art in itself, and simple approaches like comparing stack traces or the paths taken through software are insufficiently precise. The only accurate way to estimate the number (and severity) of errors is to manually look for the defect in the source code and compare the logic errors leading to the different findings. [12] This is a lengthy task and requires intimate knowledge of the software under test and was thus declared out of scope for this project.

4.1 Comparison to the Original Paper

Figure 2 is taken from the original KLEE paper [1] and shows the ECDF of the coverage measured in their experiments. Figure 3 shows the same measurements from the four runs done with the same settings as in the original paper.

It is important to note that the settings proposed by the documentation [7] do not include failing system calls, thus need to be compared to “Base” (white) elements of the graph in Figure 2 only.

In general, Figures 2 and 3 do look reasonably similar, validating the approach taken in this paper.

However, there clearly is a difference visible between the two graphs, namely that the results I was able to achieve lag behind those reported in the original paper. Since discrete numbers from the original paper unfortunately are not available, and guessing results from a graph with mediocre resolution is error-prone, further numerical analysis of this difference is not performed. From looking at the variance in the four runs in Figure 3, it seems unlikely that the difference is purely statistical.

I do not have a definitive answer for the discrepancies, but the explanation might include differences in measurements, specifically in the calculation of executable lines of code (as discussed in Section 3.5), changes in the version of KLEE the original experiments were performed with and the current version of KLEE, and performance differences between the machines the experiments were run on. The latter would either require that the performance penalty incurred by using Docker is greater than the speedup gained by 15 years of hardware developments or be based on bottlenecks during the experiments, such as I/O bandwidth.




FIGURE 2: Coverage according to the original KLEE paper [1]




FIGURE 3: Coverage measured by `gcov` across four runs

4.2 Influence of Timeout

The performance drawbacks resp. benefits of my testing setup can be indirectly measured by changing the timeout passed to KLEE. This experiment was further inspired by the seminal work by Klees *et al.*, who collected a series of guidelines that should be followed to accurately measure the performance of a fuzzer. They propose that multiple runs should be performed to increase accuracy and found that “longer timeouts may be needed to paint a complete picture of an algorithm’s performance” [12].

I chose three additional durations to conduct experiments at, with a large enough difference to ensure a complete picture of KLEE’s performance across durations: 10 minutes, 6 hours and 24 hours.

4.2.1 Changes in Coverage

Figure 4 shows the ECDF of at least three runs across the different timeouts. Predictively, the coverage increases with increased timeouts. The difference between the timeouts does not seem too big, and this is confirmed by Figure 5.




FIGURE 4: Coverage measured by gcov across different timeouts




FIGURE 5: Spread of coverage by run, normalized to the mean of the util at a timeout of 60 minutes, across different timeouts

While the average of the coverage measured by `gcov` increases slightly as the timeout increases, this is not necessarily true for individual utils. Their behavior regarding the number of distinct result values they produce can broadly be grouped as follows:

- 34 utils showed a fixed, constant coverage across all runs at all timeouts.
- 10 utils showed four or fewer discrete values that the measurements jumped between.
- The remaining utils showed a result on a more continuous scale.

The trends they show across different timeouts vary wildly. The run results of 12 utils contain values that are lower than any value measured with a lower timeout. The same is true of 23 utils when it comes to the averages.




FIGURE 6: Coverage across times on `csplit`




FIGURE 7: Coverage across times on `expand`




FIGURE 8: Coverage across times on `du`




FIGURE 9: Coverage across times on `chown`




FIGURE 10: Coverage across times on `dircolors`




FIGURE 11: Coverage across times on `mkfifo`

I selected a few utils to discuss here to give an introduction on the kinds of different outcomes produced.

- Figure 6 represents the norm and what would be expected. With increased analysis time, the minimum, maximum, and average of the results increase.
- Figure 7 is an example of what I called discrete values above. It still represents a strict increase in value across times.




FIGURE 12: Coverage across times on `sleep`




FIGURE 13: Coverage across times on `sort`




FIGURE 14: Coverage across times on `tac`

- Figure 8 shows an effect that can be observed for multiple measurements and multiple utils: As the time increases, a threshold is reached and further increasing time no longer increases the recorded value.
- Figure 9 represents an util from where the minimum value no longer strictly increases. This is likely due to statistical randomness (although this can not be said definitively without additional data points).
- Figure 10 and Figure 11 show two utils that produced unexpected outliers. While a single run for `dircolors` with a timeout of 24 hours suddenly produces only half the coverage, a single run with just 10 minutes on `mkfifo` managed to cover approximately two thirds of the missing code. The latter can be explained by the randomness inherent to KLEE's input selector (refer to Section 2.2).
- Figure 12 shows a counterintuitive result: With increased time, the reached coverage seems to decrease. I do not have an explanation for this effect, but additional data points might reduce the likeliness of this appearing due to statistical randomness.
- Figure 13 seems to no longer show any dependency between lines covered by KLEE's analysis and the timeout passed to KLEE. Just as above, this might be due to statistical randomness.
- Figure 14 finally shows a similar seeming non-dependence between time and coverage, but with discrete values and a seeming approach to a number different from the maximum. This again might be due to statistical randomness.

4.2.2 Changes in Number of Errors

Looking at changes in the number of errors in Figure 15 paints an incomplete picture: While the number of errors increases with the additional analysis time, looking at the changes in individual utils shows that there exist no utils where no issues were found with more than one hour of analysis but where lower timeouts produced findings.

The graphs for the number of each error across time for each util can be found in the repository [11]. Their behavior can be categorized as follows:

- KLEE found a constant number of errors in most utils.
 - For most that number was zero.
 - In `cat`, `csplit`, `dd`, `md5sum`, `mkdir`, `mkfifo`, `mknod`, `od`, `seq`, `split`, and `unexpand` KLEE found a single pointer error in all runs.

- In `ls`, it found a single error of type `exec` in all runs.
- For 19 utils, strict increases in number of errors across utils were recorded. This means that the number of errors found at a timeout is equal or larger compared to what was found at a shorter timeout. As an example, the average number of total errors found in the util `ptx` in one hour was 5. However, the runs with a timeout of 24 hours produced 16, 16, and 21 errors.
- The three remaining utils show a different graph:
 - `expr` produced a query timeout (so an error of type `solver`) in one run with a timeout of one hour only. Other than that, the number of total errors found was a constant zero.
 - `shred` continually produced one error of type `exec`, a strict increase between zero and one errors of type `model` and `pointer` each, and a strict increase for errors of type `solver`. However, because of the distribution of those across runs, the total is not a strict increase, because the only run to only have produced one error in total has a timeout of one hour.
 - `chgrp` only showed errors of type `model`. All four hours with a timeout of one hour produced the error, but only two out of three runs with a timeout of six hours and one out of three with a timeout of 24 hours showed the same error.




FIGURE 15: Spread of total number of errors found by run, normalized to the mean of the util at a timeout of 60 minutes, across different timeouts

5 Testing More Recent Versions of coreutils

As a second part, I wanted to see how the results from performing the same experiments on different versions of `coreutils` would behave. I chose two additional versions: 9.4, which is the current version at the time of the experiments, and 8.25, which was released exactly halfway between 6.10 and 9.4.

5.1 Differences in Testing Setup

The fundamental approach for the newer versions of `coreutils` remained the same: use old compilers on old systems where necessary and copy the binaries to KLEE’s Docker image. Specifically, the following changes compared to the Dockerfile described in Section 3 were necessary:

- For coreutils 8.25 Ubuntu 16.04 was chosen as the base image according to the same logic described in Section 3.4. Similarly, Ubuntu 22.04 was chosen for coreutils 9.4, since KLEE requires at most LLVM 13 and Ubuntu 24.04 (which is the latest available LTS version available) no longer supports a version of LLVM compatible with KLEE.
- The step replacing parts of the source code of `sort.c`, as described in [7], was changed from replacing $(128 * 1024)$ with (1024) (compared to from $(1024 * 1024)$ to 1024 in coreutils 6.10) in both coreutils 8.25 and 9.4, since the source code was changed in the meantime.
- For coreutils 8.25, LLVM and clang need to be installed at version 3.5, for coreutils at version 13.
- For coreutils 8.25, pip needed to be upgraded to version 19, because otherwise a python incompatibility prevented the installation of WLLVM.
- For coreutils 8.25, the WLLVM version remained unchanged compared to 6.10, but version 1.3.1 was chosen for coreutils 9.4.
- The `configure` step when building both coreutils 8.25 and 9.4 requires an environment variable to run as root (which is the default user of Ubuntu’s Docker images).
- Finally, the optimization flags for the LLVM bytecode generation step needed to be changed. According to the documentation, while just using `-O0` on its own is fine for LLVM 3.4 (which is what is used for coreutils 6.10), this is no longer recommended for more recent versions. [8], [13] However, I could not get the compiler to build coreutils 8.25 using the proposed solution of `-O1 -Xclang -disable-llvm-passes`. Specifically, I needed to remove the `-Xclang` argument to get the build stage to complete successfully.

5.2 Findings

I chose a fixed timeout of 60 minutes as a constant to compare the different versions of coreutils. I then executed three runs of each additional version.

5.2.1 Changes in Coverage

Figure 16 shows the coverage changes measured. The results indicate that the variance compared to version 6.10 predictively increases between the versions, as the code incrementally changes. Certain utils seem to become better penetrable by KLEE, while the coverage of others decreases, at times catastrophically. The overall average decreases across versions, which could be explained by increased complexity in the software as features are added.

5.2.2 Changes in Number of Errors

The number of errors found by KLEE remains fairly constant across the software versions and seems to even increase slightly for version 9.4. This is surprising to me, as the number of software errors should ideally decrease with time as bugs are fixed. However, this effect might be mitigated by new code adding new features or adapting the software to a changing environment introducing new errors. It might also show that running off-the-shelf fuzzing tools can be a path to finding undiscovered errors in software systems.

However, this speculation is difficult to maintain in the face of my inability to crash any util by manually replaying the inputs created by `ktest-tool` for a sample of errors reported on version 9.4. This would indicate that many (if not all) of the reported errors are in fact false positives. However, to draw definitive conclusions about this, one would need to go through every error reported and manually check them. And, as argued in Section 4, this could not be done as part of this project.




FIGURE 16: Spread of coverage by run, normalized to the mean of the util at version 6.10, across different versions




FIGURE 17: Spread of total number of errors found by run, normalized to the mean of the util at version 6.10, across different versions

6 Discussion

In this project, I successfully recreated the experiments outlined in the seminal work by Cadar *et al.* in 2008. I presented the project setup and specific results along possible interpretations. In this section, I will discuss the results in a greater context, what questions could be answered and what work remains.

6.1 Research Questions

Section 1 introduced a series of research questions to be answered in this report. Here, I summarize the answers discussed in Sections 3, 4, and 5.

Section 3 showed that it is possible to run the current version of KLEE on coreutils 6.10, even though it requires multiple assembly steps across different Docker images to resolve version incompatibilities. The original paper provided two forms of measurements: coverage as reported by `gcov` on executable lines of code and number of errors. Examining the specific errors found by KLEE is too complex and big a task to be included in this project. In the face of this, I opted to using the number of errors as an easier to calculate if inaccurate proxy. With additional steps during each run, `gcov` can be employed to measure coverage. Since no documentation is available on how exactly the authors of the original paper measured executable lines of code, comparing these measurements might not be valid.

Section 4 goes into further detail on the variance between test runs. It then continues by contrasting the results reported by the original KLEE paper with the results measured in the experiments done during this project. For code coverage, my test results were similar to what was reported by Cadar *et al.*, validating the experiments. However, there still remain unexplained minor differences. Finally, the influence of different timeouts on the results is examined. Additional analysis time predictively leads to an increased coverage and number of errors found, however certain utils behave unpredictably when looked at individually. Additional experiments might reduce statistical randomness and thus create better clarity.

Section 5 discusses the results of the same experiment on different versions of coreutils. It introduces the changes necessary to run the test suite built in this project on versions 8.25 and 9.4. The results show a slightly decreasing code coverage with increased version number and a steady number of errors found. However, it also discusses how those errors might be largely false positives.

6.2 Produced Artifacts

Since all code used for this project is publicly available [11], it can be used for further experiments. Specifically, the following artifacts have been created:

1. A framework for running KLEE-based analysis on coreutils. This includes:
 - Dockerfiles for three different versions of coreutils (6.10, 8.25, and 9.4), which document the steps necessary to build all required artifacts to run KLEE and gather extensive metrics.
 - `analyze.sh`, a shell script responsible for the actual analysis, including arguments assembly, input and output handling, and metrics collection.
 - `run-suite.py`, a python script allowing to run experiments on all (or only a subset of) coreutils efficiently, including handling of arguments and output files.
2. Extensive scripts analyzing the results of the above framework in `analyze.ipynb` and the corresponding plots:
 - For each measurement spread (like Figures 5, 15, 16, and 17) and ECDF (like Figures 3 and 4) plots, displaying the development of the value over KLEE timeouts and coreutils versions.
 - For each measurement and util plots displaying the results (similar to Figures 6, 7, 8, 9, 10, 11, 12, 13, and 14).

6.3 Future Work

Future work can be broadly categorized into two three parts:

First, additional runs with the same setup will improve the reliability of the results measured by reducing statistical randomness. This is further discussed in Section 4.

Second, the results could be improved by ensuring that both the project setup (such as compilation arguments) and the way measurements are done are exactly the same as in the original paper. The latter is especially important since the way coverage is measured with `gcov` is not described precisely enough in the original paper (refer to Section 3.5 for details).

Finally, the errors found in the different runs need to be examined further to allow further validation of the experiment setup by checking if the same software defects were found by the original authors and in my experiments. Further, it might provide additional insight into the comparison of different coreutils versions. To start, one might analyze two consecutive versions and check if bugs mentioned in the patch notes can be found in KLEE's output.

Bibliography

- [1] C. Cadar, D. Dunbar, and D. Engler, “Klee: Unassisted and automatic generation of high-coverage tests for complex systems programs,” in *Proceedings of the 8th USENIX Conference on Operating Systems Design and Implementation*, ser. OSDI’08, San Diego, California: USENIX Association, 2008, pp. 209–224.
- [2] “KLEE symbolic execution engine.” (2024), [Online]. Available: <https://klee.github.io> (visited on Jan. 24, 2024).
- [3] V. Huber, “Challenges and mitigation strategies in symbolic execution based fuzzing through the lens of survey papers,” Dec. 2023. [Online]. Available: <https://github.com/riesentoaster/review-symbolic-execution-in-fuzzing/releases/download/v1.0/Huber-Valentin-Challenges-and-Mitigation-Strategies-in-Symbolic-Execution-Based-Fuzzing-Through-the-Lens-of-Survey-Papers.pdf>.
- [4] “Institute of applied information technology.” (2024), [Online]. Available: <https://www.zhaw.ch/en/engineering/institutes-centres/init/> (visited on Jan. 24, 2024).
- [5] S. Flum and V. Huber, “Ghidrion: A ghidra plugin to support symbolic execution,” Bachelor’s Thesis, Zürich University of Applied Science — Institute of Applied Information Technology, Jun. 2023. [Online]. Available: <https://valentinhuber.me/assets/ghidrion.pdf>.
- [6] “The simple theorem prover.” (2024), [Online]. Available: <http://stp.github.io> (visited on Jan. 24, 2024).
- [7] “OSDI’08 coreutils experiments.” (2024), [Online]. Available: <https://klee.github.io/docs/coreutils-experiments/> (visited on Jan. 24, 2024).
- [8] “Tutorial on how to use KLEE to test GNU coreutils.” (2024), [Online]. Available: <https://klee.github.io/tutorials/testing-coreutils/> (visited on Jan. 24, 2024).
- [9] Y. Li, S. Ji, Y. Chen, *et al.*, “UNIFUZZ: A holistic and pragmatic Metrics-Driven platform for evaluating fuzzers,” in *30th USENIX Security Symposium (USENIX Security 21)*, USENIX Association, Aug. 2021, pp. 2777–2794, ISBN: 978-1-939133-24-3. [Online]. Available: <https://www.usenix.org/conference/usenixsecurity21/presentation/li-yuwei>.
- [10] “Whole program llvm.” (2024), [Online]. Available: <https://github.com/travitch/whole-program-llvm> (visited on Jan. 24, 2024).
- [11] “Repository of this project.” (2024), [Online]. Available: <https://github.com/riesentoaster/klee-coreutils-experiments> (visited on Jan. 24, 2024).
- [12] G. Klees, A. Ruef, B. Cooper, S. Wei, and M. Hicks, “Evaluating fuzz testing,” in *Proceedings of the 2018 ACM SIGSAC Conference on Computer and Communications Security*, ser. CCS ’18, Toronto, Canada: Association for Computing Machinery, 2018, pp. 2123–2138, ISBN: 9781450356930. DOI: 10.1145/3243734.3243804. [Online]. Available: <https://doi.org/10.1145/3243734.3243804>.
- [13] “KLEE: -O0 is not a recommended option for clang.” (2024), [Online]. Available: <https://github.com/klee/klee/issues/902> (visited on Jan. 24, 2024).

Appendix

1 Building coreutils 6.10 on KLEE's Docker image

1.1 Error

```
1130 depbase='echo freadahead.o | sed 's|[/]*$|.deps/&|;s|\.o$||'';\
1131 gcc -I. -I../lib -g -O2 -MT freadahead.o -MD -MP -MF $depbase.Tpo -c -o
    ↳ freadahead.o ../lib/freadahead.c &&\
1132 mv -f $depbase.Tpo $depbase.Po
1133 ../lib/freadahead.c: In function 'freadahead':
1134 ../lib/freadahead.c:64:3: error: #error "Please port gnulib freadahead.c to your
    ↳ platform! Look at the definition of fflush, fread on your system, then report
    ↳ this to bug-gnulib."
1135 64 | #error "Please port gnulib freadahead.c to your platform! Look at the
    ↳ definition of fflush, fread on your system, then report this to bug-gnulib."
1136 | ~~~~~
1137 make[2]: *** [Makefile:1245: freadahead.o] Error 1
1138 make[2]: Leaving directory '/home/klee/coreutils-6.10/obj-llvm/lib'
1139 make[1]: *** [Makefile:905: all] Error 2
1140 make[1]: Leaving directory '/home/klee/coreutils-6.10/obj-llvm/lib'
1141 make: *** [Makefile:769: all-recursive] Error 1
```

1.2 freadahead.c

```
1 /* Retrieve information about a FILE stream.
2    Copyright (C) 2007 Free Software Foundation, Inc.
3
4    This program is free software: you can redistribute it and/or modify
5    it under the terms of the GNU General Public License as published by
6    the Free Software Foundation; either version 3 of the License, or
7    (at your option) any later version.
8
9    This program is distributed in the hope that it will be useful,
10   but WITHOUT ANY WARRANTY; without even the implied warranty of
11   MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
12   GNU General Public License for more details.
13
14   You should have received a copy of the GNU General Public License
15   along with this program. If not, see <http://www.gnu.org/licenses/>. */
16
17 #include <config.h>
18
19 /* Specification. */
20 #include "freadahead.h"
21
22 size_t
23 freadahead (FILE *fp)
24 {
25   #if defined _IO_ferror_unlocked /* GNU libc, BeOS */
26     if (fp->_IO_write_ptr > fp->_IO_write_base)
27       return 0;
28     return fp->_IO_read_end - fp->_IO_read_ptr;
29   #elif defined __sferror /* FreeBSD, NetBSD, OpenBSD, MacOS X, Cygwin */
30     if ((fp->_flags & __SWR) != 0 || fp->_r < 0)
31       return 0;
32     return fp->_r;
33   #elif defined _IOERR /* AIX, HP-UX, IRIX, OSF/1, Solaris, mingw */
34     # if defined __sun && defined _LP64 /* Solaris/{SPARC,AMD64} 64-bit */
35     #   define fp_ ((struct { unsigned char *_ptr; \
36                             unsigned char *_base; \
37                             unsigned char *_end; \
38                             long _cnt; \
39                             int _file; \
40                             unsigned int _flag; \
```

```

41         } *) fp)
42     if ((fp->_flag & _IOWRT) != 0)
43         return 0;
44     return fp->_cnt;
45 # else
46     if ((fp->_flag & _IOWRT) != 0)
47         return 0;
48     return fp->_cnt;
49 # endif
50 #elif defined __UCLIBC__                /* uClibc */
51 # ifdef __STDIO_BUFFERS
52     if (fp->__modeflags & __FLAG_WRITING)
53         return 0;
54     return fp->__bufread - fp->__bufpos;
55 # else
56     return 0;
57 # endif
58 #elif defined __QNX__                  /* QNX */
59     if ((fp->_Mode & 0x2000 /* _MWRITE */) != 0)
60         return 0;
61     /* fp->_Buf <= fp->_Next <= fp->_Rend */
62     return fp->_Rend - fp->_Next;
63 #else
64     #error "Please port gnuilib freadahead.c to your platform! Look at the definition of
        ↪ fflush, fread on your system, then report this to bug-gnulib."
65 #endif
66 }

```

2 Spread Plots of All Measurements

Plots for individual utils are available in the project repository [11].

2.1 By Timeout




FIGURE 18: Spread of coverage as measured by gcov by run, normalized to the mean of the util at a timeout of 60 minutes, across different timeouts




FIGURE 19: Spread of branch coverage as measured by KLEE by run, normalized to the mean of the util at a timeout of 60 minutes, across different timeouts




FIGURE 20: Spread of instruction coverage as measured by KLEE by run, normalized to the mean of the util at a timeout of 60 minutes, across different timeouts






FIGURE 21: Spread of number of errors of type abort by run, normalized to the mean of the util at a timeout of 60 minutes, across different timeouts






FIGURE 24: Spread of number of errors of type ptr by run, normalized to the mean of the util at a timeout of 60 minutes, across different timeouts




FIGURE 25: Spread of number of errors of type solver by run, normalized to the mean of the util at a timeout of 60 minutes, across different timeouts




FIGURE 26: Spread of total number of errors by run, normalized to the mean of the util at a timeout of 60 minutes, across different timeouts

2.2 By Version




FIGURE 27: Spread of coverage as measured by `gcov` by run, normalized to the mean of the util at version 6.10, across different versions




FIGURE 28: Spread of branch coverage as measured by `KLEE` by run, normalized to the mean of the util at version 6.10, across different versions




FIGURE 29: Spread of instruction coverage as measured by KLEE by run, normalized to the mean of the util at version 6.10, across different versions




FIGURE 30: Spread of number of errors of type abort by run, normalized to the mean of the util at version 6.10, across different versions




FIGURE 31: Spread of number of errors of type exec by run, normalized to the mean of the util at version 6.10, across different versions




FIGURE 32: Spread of number of errors of type model by run, normalized to the mean of the util at version 6.10, across different versions




FIGURE 33: Spread of number of errors of type ptr by run, normalized to the mean of the util at version 6.10, across different versions




FIGURE 34: Spread of number of errors of type solver by run, normalized to the mean of the util at version 6.10, across different versions




FIGURE 35: Spread of total number of errors by run, normalized to the mean of the util at version 6.10, across different versions