

Running KLEE on GNU coreutils

Valentin Huber

Institute of Applied Information Technology
Zürich University of Applied Sciences ZHAW
contact@valentinhuber.me

January 24, 2024

Contents

1	Introduction	1
2	Reproducing the Original Paper	2
2.1	Project Setup	2
2.1.1	Building coreutils 6.10 on a Current Version of Ubuntu . .	2
2.1.2	Using an Old Version of Ubuntu	2
2.1.3	LLVM	2
2.1.4	Running KLEE	2
2.2	Analyzing the Results	2
2.2.1	Gathered Metrics	2
2.2.2	Comparison to the Original Paper	2
3	Analysis of the Results	2
3.1	Metrics Distribution	2
3.2	Influence of Testing Timeout	2
4	Testing More Recent Versions of core- utils	2
4.1	Differences in Testing Setup	2
4.2	Findings	2
5	Discussion	2
5.1	Research Questions	2
5.2	Produced Artifacts	2
5.3	Future Work	2
	Bibliography	2

1 Introduction

KLEE [1] is an open source, symbolic execution based, advanced fuzzing platform. It was introduced in the seminal paper titled “KLEE: Unassisted and Automatic Generation of High-Coverage Tests for Complex Systems Programs” in 2008. In their article, Cadar

et al. present their work and evaluate it on a diverse set of programs. The most prominent of those is the GNU coreutils suite, in which ten fatal errors were found.

Ever since then, KLEE has not only matured as a fuzzer, it has also been used extensively as a platform for other researchers to build on top of, as I have discovered in [3]. As an introduction to the practical side of fuzzing, I attempted to answer the following questions about KLEE:

1. Reproducing the original paper (see Section 2)
 - (a) Can the current version of KLEE be run on coreutils version 6.10 (as tested in the original paper)?
 - (b) Can the same metrics as measured in the original paper still be measured?
 - (c) How do the measured metrics compare to what was published 15 years ago?
2. Examining the statistical distribution of results over different fuzzing times (see Section 3)
 - (a) How does the non-determinism in KLEE influence the variance in the results between different test runs?
 - (b) How do different testing timeouts influence results?
3. Testing more recent versions of coreutils (see Section 4)
 - (a) What needs to change in the test setup to test more recent versions of coreutils?
 - (b) How do the results from testing different versions of coreutils differ?

All experiments were run on a virtualized server with the following specs: AMD EPYC 7713 64C 225W 2.0GHz Processor, 1 TiB RAM, 2x 25GiB/s Ethernet.

2 Reproducing the Original Paper

I'm basing my experiment setup on the original paper [2], the FAQs in the project's documentaiton [4] and the tutorial on testing coreutils version 6.11 [5].

2.1 Project Setup

2.1.1 Building coreutils 6.10 on a Current Version of Ubuntu

2.1.2 Using an Old Version of Ubuntu

2.1.3 LLVM

2.1.4 Running KLEE

2.2 Analyzing the Results

2.2.1 Gathered Metrics

2.2.2 Comparison to the Original Paper

3 Analysis of the Results

3.1 Metrics Distribution

3.2 Influence of Testing Timeout

4 Testing More Recent Versions of coreutils

4.1 Differences in Testing Setup

4.2 Findings

5 Discussion

5.1 Research Questions

5.2 Produced Artifacts

5.3 Future Work

Bibliography

[1] "KLEE symbolic execution engine." (2024), [Online]. Available: <https://klee.github.io> (visited on Jan. 24, 2024).

[2] C. Cadar, D. Dunbar, and D. Engler, "Klee: Unassisted and automatic generation of high-coverage tests for complex systems programs," in *Proceedings of the 8th USENIX Conference on Operating Systems Design and Implementation*, ser. OSDI'08, San Diego, California: USENIX Association, 2008, pp. 209–224.

[3] V. Huber, "Challenges and mitigation strategies in symbolic execution based fuzzing through the lens of survey papers," Dec. 2023.

[4] "OSDI'08 coreutils experiments." (2024), [Online]. Available: <https://klee.github.io/docs/coreutils-experiments/> (visited on Jan. 24, 2024).

[5] "Tutorial on how to use KLEE to test GNU coreutils." (2024), [Online]. Available: <https://klee.github.io/tutorials/testing-coreutils/> (visited on Jan. 24, 2024).